

Embedded Operating Systems and Architectures notes

Luca Seggiani
2026

1 Lesson 23-09-26

1.1 Introduction

Computing systems are everywhere: cars, trains, aircraft, satellites, medical devices, printers, phones, industrial controllers. In an **embedded** system, we consider *software plus hardware*, not software alone: the program only makes sense together with the machine it drives. Complexity is growing dramatically. A modern car carries ≈ 100 electronic control units and on the order of 10^8 lines of on-board code, against a handful of hundreds of lines in the 1970s.

1.1.1 Operating systems

The natural question to ask is *why* an embedded system might need an operating system. The answers are many:

- Bare metal applications no longer cope with the complexity even in the embedded systems domain;
- **An operating system raises the abstraction level:** it gives application code a small set of rules and services instead of ad-hoc hardware handling;

Our goal is to connect three things that are usually taught apart: **computer architecture**, **operating systems**, and the **computer system** that has to run both.

1.1.2 Topics

The main roadmap of the course is:

- We will connect operating-system concepts to executable embedded software;
- This course lives in the **platform** layer of a computing node: from the hardware up to the system-call interface, with the application used as the thing that exercises it.

The covered topics are:

- **Minimal computing-system architecture** — juste enough to write operating systems code;
- **Bare-metal development flow** — cross-compilation, emulators and simulators, debugging and profiling, virtual platforms;
- **Operating systems for embedded systems** — real-time and process scheduling, the anatomy of an embedded OS, architecture and services of real systems (FreeRTOS);
- **Programming an embedded OS** — building it, device management and drivers, drivers for custom devices.

1.2 Introduction to embedded operating systems

An **OS** (*Operating System*) is nothing but a program, with a very simple purpose: to mediate between hardware and the programs that use it.

There are three classic ways to draw that boundary:

- **Monolithic:** kernel is one large program; drivers, filesystems, scheduling all share one address space. Both Linux and Windows fit this description.

Monolithic kernels are simple to implement once the task switching machinery is running: once we can switch to kernel mode, we can touch any resource available to the system to carry out tasks;

- **Microkernel:** kernel does the bare minimum (IPC, scheduling, basic memory); everything else runs as separate services.

Microkernels are interesting because they are inherently safer: the service handling the filesystem, for instance, cannot corrupt other services on failure (as it runs itself as a protected service with no access to other resources). The problem comes up when complex tasks have to be performed: the main overhead is given by the IPC and scheduling facilities which sit in between services;

- **Flat:** no kernel/user separation at all; a small set of statically-defined tasks share the CPU under one scheduler.

This is important to us because embedded systems often don't feature protection systems, and as we know complete protection is impossible without hardware support.

The core of the operating system is basically the *scheduler*: everything else can be seen as a task either the operating system or the user programs carry out, based on the resources allocated by the scheduler.

1.2.1 Hardware resources

The purpose of an OS is to make hardware *usable*. Without an OS, every application must decide how to initialise the CPU, program a timer, share a UART, and recover from an interrupt. An OS supplies a small set of rules and mechanisms so application code can ask for useful services instead of repeatedly rebuilding hardware policy.

The resources are finite, especially on a small board:

- **CPU time:** which ready task runs now? We also have to take into account the fact that CPUs interact with many other devices on the same board;
- **memory:** which code and data occupy RAM and Flash?
- **devices:** who may configure a timer or consume a received byte?
- **events:** how does an interrupt become work for the right task?

1.2.2 Types of operating systems

The three labels we gave earlier describe where operating-system services execute and how strongly they are isolated.

Type	Services offered	Tradeoff
Monolithic	Together in kernel space	Direct calls, large trusted core
Microkernel	Mostly as separate servers	Isolation, IPC overhead/design
Flat	Together in one firmware image	Tiny and direct, little isolation

Monolithic A monolithic kernel puts scheduling, drivers, filesystems, and networking in one kernel address space. Services can call each other directly and quickly.

- Linux is the familiar large-scale example
- A faulty driver can corrupt kernel state

- The design is practical when the feature set is broad and tightly coupled

On a protected machine, applications normally execute outside this trusted kernel address space.

Microkernel A microkernel retains only the most fundamental mechanisms in the kernel, usually scheduling, address-space management, and inter-process communication.

Drivers and other services become separate programs that exchange messages.

- Smaller trusted computing base
- A service failure can be contained or restarted
- More interfaces and message paths to design correctly

seL4 and QNX are useful reference points for this family.

Flat Flat embedded firmware usually has one linked image and one physical address space. Tasks cooperate through shared memory and the scheduler.

There is no process loader, no virtual-memory translation, and usually no hardware-enforced boundary between tasks. This puts more pressure on software developers, as user code must be free of bugs which could potentially crash the whole operating system. This is particularly critical in systems like vehicles, medical or industrial. We know that monolithic operating systems give us an extra layer of safety, by switching to a safer kernel mode to perform critical tasks, which user code can't normally access.

1.2.3 Operating system roadmap

Where does this course sit in this hierarchy?

- **Linux** (a monolithic kernel) and *seL4/QNX* (microkernels) are the extremes we'll read about;
- This course builds **flat** first (*scratch-os*) — the simplest point on the spectrum, and the one that makes the boot-to-scheduling pipeline fully visible. *scratch-os* is intentionally small. Its value is that every important step can fit in mind;
- Then **FreeRTOS** — closer to a minimal microkernel-flavored RTOS: a small kernel, real scheduling policies, but still no memory protection between tasks by default.

FreeRTOS provides established implementations of the mechanisms we first meet in *scratch-os*:

- Task creation and priorities;
- Delay and timer-driven scheduling;
- Queues, semaphores, and mutexes;
- Interrupt-to-task communication.

The comparison matters: an API call becomes easier to reason about when its underlying state transitions are no longer invisible.

1.2.4 Real-time operating systems

The **real-time** label applied to an OS is about *bounded response*, not raw speed. The idea is that we can know exactly how long a task will take before executing it. A result computed after its deadline can be a failure even when its value is otherwise correct. This is particularly useful in embedded systems which might have external (e.g. physical) constraints, in addition to technical ones.

Requirement	Example
Functional	sample a sensor correctly
Temporal	sample it every 1 ms, before the deadline
Robustness	continue safely when one input is missing

1.2.5 Embedded operating systems

Embedded operating systems usually look different than traditional ones:

- **Resource-constrained:** kilobytes of RAM/flash, not gigabytes — no room for a general-purpose kernel's overhead;
- **Real-time:** correctness often depends on *when* something happens, not just *whether* it happens;

- **No user protection needed (usually):** one vendor, one firmware image, one set of tasks — the isolation problems a desktop OS solves mostly don't apply here.

This is mainly a cost problem: it's not that car manufacturers don't *want* safety in their devices, but that CPUs which don't feature protection mechanisms are usually cheaper;

- These constraints are why *scratch-os* can be simple and still be a real, working OS.

An embedded OS is selected and configured around a particular product, not an unknown future workload.

- Static task counts and fixed buffers can replace dynamic allocation;
- A simple priority scheduler can replace a general-purpose fairness policy;
- Direct register access can replace a layered device stack;
- Predictable execution is often more valuable than maximum throughput.

Small is not incomplete when it matches the system's responsibilities.

Many **MCU** (*Micro Controller Unit*) products ship one firmware image that includes startup code, drivers, the OS, and application logic. Updating that image changes the whole system.

This simplifies deployment, but it raises the bar for discipline:

- a bad pointer can damage any writable state
- a blocked high-priority task can stop an important service
- recovery and observability must be designed into the firmware

Later, we will see how scheduling and synchronisation reduce these risks.

1.2.6 Embedded systems

There are many ways we can classify **embedded systems**:

- A computer built into a larger device to do **one dedicated job** — not a general-purpose machine you install arbitrary software on;
- Hardware + software (often called *firmware*) shipped as one unit; the end user never sees a shell or a package manager;
- Everywhere: engine controllers, pacemakers, drones, washing machines, Wi-Fi routers, the microcontroller in a USB charger;
- The population of embedded processors passed the human population decades ago — most CPUs made are never in a "computer".

We can compare embedded systems to general-purpose ones (such as modern computer or phone operating systems):

	Embedded	General-purpose
Workload	one fixed task	anything the user installs
Resources	KB of RAM/flash	GB of RAM, TB of disk
Timing	often hard deadlines	best-effort
User model	no interactive user	multi-user, multi-tasking
Cost/power	pennies, milliwatts	dollars, tens of watts

The right-hand column is what a desktop OS is built for; scratch-os targets the left-hand one, which is *why* it can be small.

1.3 Microcontrollers

We begin to approach the device that we'll actually need to program. A **microcontroller** or **MCU** (*Micro Controller Unit*) integrates onto a single chip what a desktop splits across a whole motherboard:

- **CPU core** — executes instructions (ours: ARM Cortex-M3);
- **RAM** — volatile working memory, tens of KB;
- **Flash** — non-volatile program storage; code runs directly from it;
- **Peripherals** — timers, UART, GPIO, ADC, all on-chip;
- One chip, one clock, one address space — no external memory bus to wire up, no BIOS, no disk.

Compared to desktop or handled general-purpose CPUs, we note that:

- A desktop CPU is *just* the core: RAM, storage, and I/O are separate chips reached over buses; an MMU gives every process its own virtual address space;
- An MCU has **one flat physical address map** shared by code, data, and every peripheral register — no MMU, no virtual memory (a Cortex-M has at most an optional MPU);
- On an MCU, “the hardware” is *memory you write to*: there is no driver layer between your code and a device register.

1.3.1 Reset routine

As an example, let's look into what happens to a Cortex-M CPU at reset time. After reset, the Cortex-M does not discover an operating system. It reads two words from the vector table in Flash:

1. The initial value for the main stack pointer
2. The address of `Reset_Handler`.

The reset handler is our first code. It establishes the C runtime environment before calling `main()`.

1.3.2 Registers

The CPU operates on values held in registers. Memory is accessed explicitly with load and store instructions.

Register role	Meaning
Program Counter	Selects the next instruction
Stack Pointer	Locates calls, local variables, saved state
General Registers / Programmer Registers	Hold operands and addresses
Status Register	Records conditions and interrupt state

When an exception occurs, the processor automatically saves essential state on the current stack so execution can later resume.

1.3.3 Memory

We can talk about the *memory* available to an MCU. Usually there are two kinds of memory:

- **SRAM / DRAM:** not persistent, but extremely fast. Used as main program memory;
- **Flash** memory: read-only, but still very fast and fit (in embedded systems) for containing code. This is in accord with the usual definition of *Harvard Architecture*. We should note that most general purpose devices use the opposite architecture, that is *Von Neumann Architecture* (which use program memory both for code and data).

Flash keeps its contents without power. SRAM is fast writable working memory whose contents disappear on reset or power loss.

```
Flash: vector table | .text | .rodata | initial bytes for .data
SRAM:  .data       | .bss  | heap area | stack
```

This split explains why startup code must copy some bytes and clear others before ordinary C code can safely run.

2 Lesson 30-09-26

2.1 Memory

Last week we went over the structure of an embedded system MCU. Now we focus on memory, with the target of defining a memory map and defining a linker script for it.

2.1.1 Types of memory

Let's quickly go over the *types* of memory available to us.

SRAM

- **SRAM** (static RAM) sits on the CPU's memory bus: every byte has an address, and the CPU reads or writes it with a normal load/store;
- **Fast:** on a microcontroller, one access takes about one clock cycle;
- **Byte-writable:** we can change a single byte, any time, as often as we like;
- **Volatile:** power off and everything is lost; at power-on the content is **undefined** (random bits, not zeros).

SRAM is where anything that **changes** while the program runs must live: variables, the stack.

NOR flash

- Also sits on the memory bus, like a ROM: every word has an address, random-access reads are fast;
- So the CPU can **fetch** instructions directly from it (common in MCUs): **execute-in-place** (XIP). No copy needed;
- **Non-volatile**: the program is still there after power-off;
- Writing is **not** a store instruction:
 - First **erase** a whole sector (e.g. a few KB) → all bits become 1;
 - Then **program** words (bits can only go 1 → 0)
 - Microseconds to milliseconds, through a special command sequence
- Result: at run time, NOR flash is effectively **read-only**.

NAND flash and disks

- NAND flash (SD cards, SSDs) and hard disks are **not** on the memory bus: they are reached through a controller, one **block** (page) at a time, e.g. 512 B – 4 KB;
- A block has no CPU address → the CPU **cannot fetch instructions** from it, nor read a single variable;
- Contents must first be **copied into RAM**, then used from there;

This is what a desktop PC does: the program sits on the SSD, and the OS **loader** copies it into RAM before running it. A microcontroller has no loader: its program already sits in NOR flash, at the address it runs from.

2.1.2 Effect on program sections

The different memory types we have available play a role in how we define our *program headers* (or *segment headers*) for embedded programs.

Technology	On the bus	Executable	Writable at runtime	Persistent
SRAM	Yes	Yes (XIP)	Yes, bytes	No
NOR flash	Yes	Yes (XIP)	No (erase/program)	Yes
NAND / disk	No (blocks)	No	Via a driver	Yes

The CPU sees one 4 GB address space. On every Cortex-M the **regions** are fixed by the architecture; each chip decides what it populates.

Address range	Region	On mps2-an385
0x0000_0000–0x1FFF_FFFF	Code	4 MB at 0x0: the program ("FLASH")
0x2000_0000–0x3FFF_FFFF	SRAM	4 MB at 0x2000_0000: "RAM"
0x4000_0000–0x5FFF_FFFF	Peripherals	UART, timers...
0x6000_0000–0xDFFF_FFFF	External RAM / devices	Not used here
0xE000_0000–0xE00F_FFFF	CPU's private registers	Control of the CPU itself (at 0xE000_E000)

Tabella 1: Cortex-M memory map and the regions populated by mps2-an385.

2.1.3 Linker script

This leaves us with the following linker script, which we can find very useful to analyze:

```

1  /* Linker script for mps2-an385 machine (ARM Cortex-M3, MPS2 platform). */
2
3  /* Physical memory regions of the board. */
4  MEMORY
5  {
6      /* code, rodata (XIP) */
7      FLASH (rx)  : ORIGIN = 0x00000000, LENGTH = 4M
8
9      /* data, bss, stack */
10     RAM   (rwx) : ORIGIN = 0x20000000, LENGTH = 4M
11 }
12
13 /* Entry point, also entry 1 of vector table (reset handler). */
14 ENTRY(Reset_Handler)
15
16 /* initial SP = top of RAM; the stack grows downward from here */
17 _estack = ORIGIN(RAM) + LENGTH(RAM);
18
19 /* 4 KB reserved for the stack*/
20 _Min_Stack_Size = 0x1000;
21
22 /* combine input (ELF) sections into output sections */
23 SECTIONS
24 {
25     /* Vector table first, at address 0, so the CPU finds the initial SP
26      * and reset vector where the architecture requires them. */
27     .isr_vector :
28     {
29         . = ALIGN(4);           /* vector table must be word-aligned */
30         KEEP(*(.isr_vector))    /* KEEP: never garbage-collect it */
31         . = ALIGN(4);           /* keep alignment */
32     } > FLASH                   /* place in FLASH */
33
34     /* Executable code and read-only data. */
35     .text :
36     {
37         . = ALIGN(4);           /* word-align the start of this section */
38         *(.text*)               /* all code from every input object */
39         *(.rodata*)             /* string literals, const tables */
40         . = ALIGN(4);           /* word-align the end so (_etext) */
41         _etext = .;             /* end-of-text marker */
42     } > FLASH                   /* place in FLASH */
43
44     /* FLASH address where .data's initial image is stored: the technique
45      allows for reserving space in RAM while storing in flash. */
46     _sidata = LOADADDR(.data);
47
48     /* Initialized variables: they live (run) in RAM but their initial
49      * contents are stored in FLASH and copied over by Reset_Handler. */
50     .data :
51     {
52         . = ALIGN(4);           /* word-align the start of .data */
53         _sidata = .;            /* start of .data in RAM */
54         *(.data*)               /* every initialized-variable input
55 section */
56         . = ALIGN(4);           /* word-align the end */
57         _edata = .;             /* end of .data in RAM - copy loop stops
58 here */
59     } > RAM AT > FLASH          /* place in FLASH, access in RAM */
60 }

```

```

59  /* Zero-initialized variables: reserve RAM, no stored image. */
60  .bss :
61  {
62      . = ALIGN(4);          /* word-align the start of .bss */
63      _sbss = .;             /* start of .bss */
64      *(.bss*)               /* every zero-initialized input section */
65      *(COMMON)              /* legacy tentative-definition symbols */
66      . = ALIGN(4);          /* word-align the end */
67      _ebss = .;             /* end of .bss - zero loop stops here */
68  } > RAM
69 }

```

For now we intentionally skip the *heap* and other advanced features, but this should be pretty useful to understand how the memory is laid out. Let us remember that here we are defining memory **regions**, declaring memory **sections**, and expecting the linker to generate for us memory **segments** (program headers) which place the data in our segments within more "blocks" (that is, segments) to load in different memory regions (the ones seen above).

2.1.4 LMA vs VMA

Let us also note the *LMA/VMA* distinction:

- **LMA** (*Load Memory Address*) is where the initial bytes of data are located in the executable/image (the program headers which we load);
- **VMA** (*Virtual/Run-time Memory Address*) is where the data in the section expects to live when the program executes.

The gimmick at lines 45-56 allows us to set a LMA for the data section inside the flash memory, and set its VMA to the address it will be copied to in RAM. Of course, the task of copying the data itself is assigned to the reset handler.

We also note that the `.bss` initialize sector has a VMA (where to access uninitialized data) but no LMA (nothing to load, the reset handler just zeroes out).

2.1.5 Vector table

One important region of the memory (more specifically, FLASH memory) is the **vector table**.

- It is a **table** of function addresses at a fixed place (address 0x0 at reset).
- Each **exception** (an event that makes the CPU stop and run a specific handler: reset, a fault, ...) has a numbered **slot**; the CPU reads slot *N* to find where to jump,
- Slot 0 is special: it contains not an address of code, but the initial stack pointer (SP), which gets loaded automatically by the CPU. This means that the reset handler is at index 1 (not 0).

In `boot/startup.s` (base), section `.isr_vector`:

```

1  isr_vector:
2      .word _estack          /* 0: initial SP (C: &_estack) */
3      .word Reset_Handler   /* 1: reset (C: Reset_Handler) */

```

The stored value of word 1 is `Reset_Handler` | 1: the low bit marks Thumb code. Nothing ever *calls* `isr_vector`; only the hardware reads it.

2.2 Registers

We now go over the ARM architecture programmer registers. In all ARM processors, the following registers are available and accessible in any processor mode:

- 13 general-purpose registers: R0–R12: user variables, with call arguments passed through R0–R3;
- One *Stack Pointer* (SP): address of the top of the stack;
- One *Link Register* (LR): where to return after a function call;
- One *Program Counter* (PC): address of the next instruction to execute;
- One *Status Register* (xPSR): flags, Thumb bit, current exception.

Some things to note at this level which may not be clear:

- There are no “variables” at this level: only these registers and memory.
- The stack grows **downwards**: a push first does $SP -= 4$, then writes.

xPSR packs three things into one register:

Bits	Content
31–28	flags N Z C V (negative, zero, carry, overflow), set by <code>cmp</code>
24	T : Thumb state — must always be 1 on Cortex-M
8–0	number of the exception being handled (0 = none)

Tabella 2: xPSR fields

The CPU is always in one of two **modes**:

- **Thread mode** — normal code: `Reset_Handler`, `main()` (xPSR bits 8–0 = 0).
- **Handler mode** — running an exception handler (bits 8–0 = its number).

Code can also be *privileged* or not; everything in this unit runs privileged.

2.3 ISA

ARM Cortex-M uses a **load/store** architecture: arithmetic and logical operations operate on registers, while memory is accessed explicitly through load and store instructions.

- `ldr`: **load** data from memory into a register;
- `str`: **store** data from a register into memory;
- Arithmetic instructions do not directly operate on memory operands.

Let’s have a table of example memory operations:

Assembly	C equivalent
<code>ldr r3, [r0]</code>	<code>r3 = *r0;</code>
<code>str r3, [r1]</code>	<code>*r1 = r3;</code>
<code>ldr r3, [r0], #4</code>	<code>r3 = *r0; r0 += 4;</code> <code>// i.e. r3 = *src++;</code>
<code>str r3, [r1], #4</code>	<code>*r1 = r3; r1 += 4;</code> <code>// i.e. *dst++ = r3;</code>
<code>strb r1, [r0, #7]</code>	<code>((uint8_t *)r0)[7] = r1;</code> <code>// one byte, at offset 7</code>

2.4 Post- and pre-increment

In ARM we have **pre-indexed** and **post-indexed** memory operations, that is control over the time when the base register is updated.

Being more specific, ARM allows load and store instructions to optionally *alter* the address contained within a base register. 3 forms are available:

- `ldr r0, [r1, #4]` is equivalent to immediate offset addressing in RISC-V, with #4 being the immediate operand;
- `ldr r0, [r1, #4]!`, with the additional `!`, means performing a *write-back* or **pre-indexing** operation. The calculated address, obtained by adding the offset, is *written back* into the base register before the load operation. The result we get is that of a load from the base address plus the offset, with the base address reflecting the change after execution. Compared to RISC-V, this saves an instruction;
- `ldr r0, [r1], #4` means performing a **post-indexing** operation. This means increasing the value of the base register by offset, after performing the load. We can think of this as being analogous to the `++` syntax available in C. This also saves an instruction compared to RISC-V.

As we can see, the main point of this complication is to save a few instructions when accessing arrays, or generally where striding memory access patterns are needed. Since a word is 4 bytes on this machine, this is useful for walking through an array or copying a block of memory.

For example:

```
1 copy:
2     ldr r3, [r0], #4
3     str r3, [r1], #4
4     b    copy
```

copies one word at a time while automatically advancing both pointers.

2.5 Loading constants and addresses

Assembly code also needs mechanisms for loading constants, referring to addresses, defining labels, and changing the control flow.

Let's have a table of example (pseudo and normal) instructions for constant loading.

Assembly	C equivalent
<code>ldr r0, =_estack</code>	<code>r0 = (uint32_t)&_estack;</code>
<code>mov sp, r0</code>	<code>sp = r0;</code>
<code>movs r3, #0</code>	<code>r3 = 0;</code>
<code>copy_data:</code>	label: a name for this address
<code>b copy_data</code>	<code>goto copy_data;</code>
<code>bl main</code>	<code>main();</code>

The pseudo-instruction

```
1 ldr r0, =_estack
```

loads the value represented by `_estack` into `r0`. This is often an address, but the syntax can be used for arbitrary 32-bit constants.

A large constant generally cannot be encoded directly in a single Thumb instruction. The assembler can therefore place the constant in a nearby **literal pool** and generate a PC-relative load.

For example, GDB may show:

```
1 ldr r0, [pc, #92]
```

The directive

```
1 .ltorg
```

asks the assembler to emit the current literal pool at that point.

Note that a label names an address in the program:

```
1 copy_data:  
2 ...
```

A branch transfers execution to that address:

```
1 b copy_data
```

This is analogous to a C `goto`, although assembly labels are addresses in the instruction stream rather than variables.

The instruction `bl` means **branch with link**. It branches to the target and simultaneously saves the return address in the **link register (LR)**. This is the normal mechanism used for a subroutine call:

```
1 bl main
```

A corresponding return normally uses the address stored in `lr`.

2.5.1 Branches

Assembly	C equivalent
<code>cmp r1, r2</code>	compute <code>r1 - r2</code> and update the flags
<code>bhs done</code>	if <code>(r1 >= r2)</code> goto <code>done</code> ; // unsigned comparison
<code>orr r1, r1, #0x10</code>	<code>r1 = 0x10</code> ;
<code>cpsid i</code>	disable interrupt acceptance
<code>cpsie i</code>	enable interrupt acceptance

Tabella 3: Conditional branches, bit operations and interrupt control.

`cmp` performs a subtraction internally (`r1 - r2`) but does not store the result. Instead, it updates the condition flags (N, Z, C, and V) in the xPSR.

A subsequent conditional branch examines these flags.

For example, `bhs` means “higher or same” for an **unsigned** comparison:

```
1 cmp r1, r2  
2 bhs done
```

This corresponds to:

```
1 if (r1 >= r2)  
2 goto done;
```

A C `while` loop therefore naturally becomes a test at the top, followed by a backwards branch at the bottom:

```
1 loop:  
2 cmp r1, r2      @ while (r1 < r2) {  
3 bhs done  
4 ...            @   body  
5 b loop          @ }  
6 done:
```

2.5.2 Interrupt masking

On Cortex-M, `cpsid i` disables the acceptance of maskable interrupts, while `cpsie i` enables them again:

```
1 cpsid i
2 ...
3 cpsie i
```

These correspond conceptually to the CMSIS functions `__disable_irq()` and `__enable_irq()`.

2.5.3 Assembler directives

Assembly source contains two fundamentally different kinds of lines:

1. **Instructions**, which become CPU instructions.
2. **Assembler directives**, which control how the assembler constructs the object file.

Directives generally begin with a period (.) and do not themselves execute on the CPU.

For example:

```
1 .section .isr_vector,"a",%progbits
2 .global isr_vector
3
4 isr_vector:
5     .word _estack
6     .word Reset_Handler
7
8 .section .text,"ax",%progbits
9 .global Reset_Handler
10 .thumb_func
11
12 Reset_Handler:
```

Here:

- `.section` selects or creates an ELF section.
- `.global` makes a symbol visible to the linker and other object files.
- `.word` emits a 32-bit value into the current section.
- `.thumb_func` tells the assembler that the following function is Thumb code.

Thus, `isr_vector` is not a function being executed by the CPU: the assembler is simply placing two 32-bit values into the `.isr_vector` section.

The corresponding idea in C is:

```
1 __attribute__((section(".isr_vector")))
2 const uint32_t isr_vector[] = {
3     (uint32_t)&_estack,
4     (uint32_t)Reset_Handler
5 };
```

Without `.global`, an assembly symbol is local to the object file, similar in effect to a static C symbol.

2.5.4 Thumb function addresses

Cortex-M executes Thumb instructions exclusively. Function pointers and exception-vector entries therefore have bit 0 set to indicate Thumb state.

For a symbol such as `Reset_Handler`, the assembler/linker can produce an address with bit 0 set when it is declared as a Thumb function:

```
1 .thumb_func
2 Reset_Handler:
```

For example, if the actual instruction address is `0x00000040`, the Thumb function pointer is represented as `0x00000041`.

The CPU uses this bit to identify the target as Thumb code; the actual instruction address is still aligned to an even address.

2.6 Faults

Let's look in more detail at **faults**, so that we can understand the whole of the *handlers* defined in the ISR vector table.

An **exception** is any event that interrupts the current code and causes the CPU to execute a handler whose address is stored in the vector table, switching to Handler mode.

#	Exception	Raised by	#	Exception	Raised by
0	<i>None</i>	Slot holds the initial SP	7–10	<i>Reserved</i>	Slot must hold 0
1	Reset	Power-on, reset button	11	SVC	The svc instruction
2	NMI	Non-maskable interrupt	12	DebugMonitor	Debug events (unused here)
3	HardFault	Faults that cannot be handled elsewhere	13	<i>Reserved</i>	Slot must hold 0
4	MemManage	Memory protection violations	14	PendSV	A software request
5	BusFault	Bus access errors	15	SysTick	Built-in timer
6	UsageFault	Invalid instruction use, division by zero, ...	16+	IRQs	Device interrupts

At reset, only **HardFault** is active among the configurable fault handlers: faults that cannot be handled by their corresponding handlers escalate to **HardFault**. This corresponds to “escalate 3” behavior observed during debugging.

2.6.1 Handlers

The initial vector table contains only two words, but the Cortex-M processor expects one word per exception vector, for a total of 16 entries in the initial table. Slots 7–10 and 13 are reserved and must contain zero. Every other slot must contain a valid handler address, even if no specific handler has been implemented yet.

A default handler can be defined as follows:

```
1 .word HardFault_Handler @ slot 3, and so on up to 15
2
3 .weak HardFault_Handler
```

```

4  .thumb_set HardFault_Handler, Default_Handler
5
6  Default_Handler:
7      b Default_Handler @ for (;;) {}

```

The equivalent declaration in C is:

```

1  void HardFault_Handler(void)
2      __attribute__((weak, alias("Default_Handler")));

```

A **weak symbol** provides a default implementation that can be overridden by a normal definition in another object file. If `main.c` defines its own `HardFault_Handler`, the linker uses that implementation instead of the default.

The completed vector table occupies 16 words:

$$16 \times 4 = 64 \text{ bytes} = 0x40.$$

Consequently, if the table starts at address `0x0`, the `Reset_Handler` entry point moves from `0x8` to `0x40`. Since Cortex-M executes Thumb instructions, its vector entry contains `0x41`, with bit 0 set to indicate Thumb state.

2.7 CPU registers

One interesting thing about ARM is that some CPU settings live in *registers* which are stored at certain **memory addresses**.

For instance, the **SCB** (*System Control Block*) is a set of 32-bit registers inside the CPU, reached like memory in the system area at `0xE000E000 – 0xE000EFFF`.

Let's recall the default processor registers:

Register	Purpose
r0–r12	General-purpose registers
r13 (SP)	Stack pointer
r14 (LR)	Link register; return address
r15 (PC)	Program counter
xPSR	Flags, execution state and exception number
MSP	Main stack pointer
PSP	Process stack pointer
PRIMASK	Masks configurable interrupts
BASEPRI	Masks interrupts below a priority threshold
FAULTMASK	Masks exceptions except NMI
CONTROL	Selects stack and privilege state in Thread mode

The exact register set depends on the Cortex-M implementation. In particular, `BASEPRI` is not implemented on every Cortex-M core.

The `SP` register refers to the currently selected stack pointer, either `MSP` or `PSP`. The processor uses the main stack pointer in Handler mode.

These peripherals expose registers through fixed memory addresses. They are accessed using ordinary load and store instructions.

Let's look at the main peripherals we have:

Block	Base address	Purpose
SCnSCB	0xE000E000	System control and feature registers
SysTick	0xE000E010	24-bit system timer
NVIC	0xE000E100	Interrupt enable, pending state and priorities
SCB	0xE000ED00	Exceptions, faults and processor configuration
MPU	0xE000ED90	Memory protection (optional)
DWT	0xE0001000	Watchpoints, cycle counter and profiling
FPB	0xE0002000	Flash patch and breakpoint unit
ITM	0xE0000000	Instrumentation and debug output
CoreDebug	0xE000EDF0	Debug control and status
TPIU	0xE0040000	Trace output interface

These addresses correspond to the Cortex-M3 memory map. Optional blocks may be absent or incompletely implemented in a particular processor or emulator.

The *System Control Block (SCB)* provides registers for processor configuration, exception handling and fault diagnosis.

Let's look at some important elements of the SCB:

Register	Offset	Purpose
CPUID	0x00	Processor identification
ICSR	0x04	Exception and interrupt state
VTOR	0x08	Vector table base address
AIRCR	0x0C	System reset and priority configuration
SCR	0x10	Sleep behavior
CCR	0x14	Processor configuration
SHPR1-3	0x18-0x23	System exception priorities
SHCSR	0x24	System handler enable and status
CFSR	0x28	Configurable fault status
HFSR	0x2C	HardFault status
DFSR	0x30	Debug fault status
MMFAR	0x34	Memory-management fault address
BFAR	0x38	Bus fault address

2.8 Real chips

QEMU's CPU is ready as soon as it leaves reset. A real MCU's startup code usually does more, often in a vendor function `SystemInit()`:

- **Clock / PLL:** the chip starts on a slow internal oscillator; the PLL multiplies the clock frequency up to the desired operating frequency.
- **Flash wait states:** Flash memory is slower than a fast CPU. The Flash controller must be configured to insert additional wait cycles *before* increasing the CPU clock frequency.
- **Watchdog:** a timer that resets the chip unless software periodically “feeds” it. On some chips the watchdog starts running immediately after reset, so the startup code must feed or disable it.
- **FPU (Cortex-M4F):** the floating-point unit is disabled at reset. It can be enabled through the CPACR register at 0xE000ED88; otherwise executing an FPU instruction can cause a UsageFault.

- **C++:** before entering `main()`, the startup code must call the static constructors stored in `.init_array`.

3 Lesson 01-10-26

3.1 Freestanding C

The C standard defines two kinds of environment:

- **Hosted:** C library (`libc`) and an OS sit under the user code, library functions such as `printf()`, `malloc()`, `fopen()` are available. `main()` is called by the `libc` startup code, called by the OS. Of course, each OS implementation implements its own `libc` (GNU/Linux, Windows, ...). There are even versions of `libc` which run on bare metal (`newlib`);
- **Freestanding:** nothing sits under the user code, no library functions are available and nobody calls `main()` (the default entry point is the reset handler).

We will be using a *freestanding* environment from now on.

System calls

In fact, `libc` in Linux is only an interface between user code and the OS: a call to a `libc` function in Linux means:

- Performing a little setup work in userland, via the `libc` implementations;
- Performing some kind of *system call*, which leaves the CPU to the kernel;
- Waiting for the kernel to carry out our system call;
- Getting the result back from `libc`.

We decide to opt out of libraries like `libc` because they still need architecture-specific hooks (`write`, `crt0`, ...) hooks to be implemented: we're better off doing everything ourselves.

Compiler helpers

One piece of help we will use is the *compiler helpers* provided by the **GCC** (*Gnu Compiler Collection*) toolchain we use. These implement some basic C operators (such as divide and modulo), which may not be implemented in hardware. Was someone to be curious, some example helpers can be found at <https://github.com/seggiani-luca/micro-sim/blob/main/eprom/src/lib/helpers.cpp> (there are likely cleaner and better implemented one somewhere in the GCC source).

3.1.1 Semihosting

We still need to be able to print stuff to console. The solution we will use is **semihosting**: the target asks the *host* (our machine) to perform the I/O.

1. The program puts an operation code in `r0` (`0x04 = SYS_WRITE0`) and the string address in `r1`. `SYS_WRITE0` is the opcode for "print this NULL-terminated string";
2. It executes `bkpt 0xab`: a breakpoint instruction with a special number;
3. QEMU (or a debugger) sees the breakpoint, reads `r0` and `r1`, reads the string from target memory, and prints it **on your laptop**;
4. The program resumes after the `bkpt`.

This only works with a debugger or emulator attached (e.g., `-semihosting` in `make run`). On a real board without a debugger, `bkpt` would simply stop the CPU.

In code

In code, this translates to:

```
1 static void semihost_write0( const char *message )
2 {
3     register uint32_t    r0 __asm__( "r0" ) = 0x04;    /* SYS_WRITE0 */
4     register const char *r1 __asm__( "r1" ) = message; /* string address
5     /*
6     __asm__ volatile ( "bkpt 0xab" : : "r"( r0 ), "r"( r1 ) : "memory" );
7 }
```

Let's look at what `bkpt 0xab` really does:

- The program **stops** doing its own work;
- Control passes to **someone more powerful** (QEMU, outside the CPU);
- That someone performs the job and **hands control back**.

This is the same shape as a system call on Linux: your program cannot drive the terminal itself, so it **asks the kernel** to do it.

Next question: can the Cortex-M3 itself have a 'powerful' part and a 'restricted' part, like kernel mode and user mode? **Yes**: privileged and unprivileged execution, next slide.

3.2 Privilege level

Even if our implementation of an OS will be *flat*, let's go over **privilege** levels in our ARM Cortex-M3 CPU.

We know that an *OS kernel* runs in kernel mode, whereas an application runs in *user mode*. Cortex-M3 has the same idea in hardware:

- **Handler mode**: used when running an exception handler, always privileged;
- **Thread mode**: normal mode used for user code, can be privileged or unprivileged.

After reset we find the CPU in thread mode, privileged. We have to manually set the CPU to unprivileged thread mode if we want protection on user code.

A special CPU register (not in memory) selects Thread-mode privilege:

Bit	Name	Meaning
0	nPRIV	0 = privileged, 1 = unprivileged (Thread mode only)
1	SPSEL	selects which stack pointer Thread mode uses (not used here)

Special registers are reached with two dedicated instructions:

Instruction	C equivalent
<code>mrs r0, control</code>	<code>r0 = CONTROL;</code> (read special register)
<code>msr control, r0</code>	<code>CONTROL = r0;</code> (write special register)

3.2.1 Dropping privilege

So, for example, to drop our privilege lever to unprivileged:

```
1 mrs r0, control    @ r0 = CONTROL
2 orr r0, r0, #1     @ r0 = r0 | 1      (set nPRIV)
3 msr control, r0    @ CONTROL = r0     -> now unprivileged
4 isb               @ discard instructions already fetched under old rules
```

Writing `CONTROL` is itself a privileged operation: an unprivileged `msr control, ...` is ignored. Code cannot promote itself back. The only ways back to privileged: reset, or an exception (its handler runs in Handler mode, always privileged).

Attempt	Result
Read/write the System Control Space (0xE000E000 ...): SysTick, NVIC, SCB (VTOR, SHCSR, CCR, priorities) <code>msr control, r0</code> (raise privilege back) <code>cpsid i / cpsie i</code> (mask interrupts) Semihosting bkpt 0xab	BusFault ignored ignored QEMU refuses it by default → fault

Ordinary RAM, Flash, and user code still work normally. The restricted code simply cannot **reconfigure** the machine or silence it.

3.2.2 Raising privilege

Unprivileged code cannot touch the hardware, but it can still *call* privileged mode handlers: `svc #n` (*SuperVisor Call*) is allowed in unprivileged code.

1. It raises the **SVCall exception** (vector-table entry 11, `SVC_Handler`). The CPU enters **Handler mode**, and the corresponding handler runs privileged;
2. When the handler returns, the CPU goes back to the instruction after `svc`, still **unprivileged**.

`#n` is a small number encoded in the instruction. Real kernels use a request number and arguments to tell which service is wanted; here we only use `svc #0`.

This allows us to implement what effectively are **system calls**.

4 Lesson 07-10-26

4.1 ARM hardware interface

The ARM hardware interface is split in 2 components:

- The **CMSDK** (*Cortex-M System Design Kit*): it is an hardware specification of ready-made peripheral designs that chip designers place next to the CPU;
- The **CMSIS** (*Cortex-M Software Interface Standard*): it is a standard for the C headers and functions that describe a chip.

Our board's peripherals are CMSDK blocks: Arm documents their registers once, the same on every chip that uses them. More specifically, it is an *MPS2*: an Arm development board built around an FPGA (a chip whose logic is loaded at power-on). *AN385* (Application Note 385) is one FPGA image for it: a Cortex-M3 plus CMSDK peripherals.

Address	Device
0x4000_0000 / 0x4000_1000 / 0x4000_1000	Timer 0, Timer 1, dual timer
0x4000_4000 / 0x4000_5000 / 0x4000_6000	UART0, UART1, UART2
0x4000_8000 / 0x4001_0000	Watchdog, GPIO pins
0x4002_8000	FPGA I/O (MPS2 block)

We use no vendor CMSIS files (no libraries at all): our register header is hand-written in CMSIS style, so its names look like those in vendor code.

Basically, we have a CMSDK peripheral set, for which we program a CMSIS interface by hand.

4.1.1 Buses

A **bus** is the set of wires (address, data, control) that carries every load and store from the CPU to a memory or a device. Arm defines two kinds:

- **AHB** (*Advanced High-performance Bus*): the main bus; pipelined, one transfer per cycle. Memories and devices that need speed sit here;
- **APB** (*Advanced Peripheral Bus*): simple, slow, low power. For registers that are touched rarely: configure a timer, send one byte

The bridge turns an AHB access into an APB access: a few extra cycles are needed for conversion. For our code, nothing changes: a register is a load or store at an address on either bus; only the timing differs.

4.1.2 External interrupts

A device can raise an interrupt request (**IRQ**) line to tell the CPU "something happened". Device line N is served by vector-table slot $16 + N$, right after the 16 core slots we saw were related to internal interrupts (exceptions, such as the Reset Handler).

IRQ	Interrupt source	IRQ	Interrupt source
0	UART 0 receive	16	GPIO 2 combined
1	UART 0 transmit	17	GPIO 3 combined
2	UART 1 receive	18	UART 3 receive
3	UART 1 transmit	19	UART 3 transmit
4	UART 2 receive	20	UART 4 receive
5	UART 2 transmit	21	UART 4 transmit
6	GPIO 0 combined	22	SPI 2
7	GPIO 1 combined	23	SPI 3, SPI 4
8	Timer 0	24	GPIO 0, pin 0
9	Timer 1	25	GPIO 0, pin 1
10	Dual Timer	26	GPIO 0, pin 2
11	SPI 0, SPI 1	27	GPIO 0, pin 3
12	UART overflow (UART 0, 1, 2)	28	GPIO 0, pin 4
13	Ethernet	29	GPIO 0, pin 5
14	Audio I2S	30	GPIO 0, pin 6
15	Touch Screen Controller	31	GPIO 0, pin 7

4.2 ARM ABI

Let's look at the ARM **ABI** (*Application Binary Interface*), as it will be useful to know how to implement interrupt handlers (we have to know how arguments are passed around).

C and assembly can call each other only if both follow the same rules: the **AAPCS** (*Procedure Call Standard for the Arm Architecture*).

Register	Role in a call	Kept across the call?
r0–r3	The first four arguments; r0 also carries the result	No: caller-saved (scratch)
r4–r11	The function's own long-lived values	Yes: callee-saved
r12	Scratch	No: caller-saved
lr	Return address, written by bl	No: a function that calls another saves it first
sp	Must be a multiple of 8 at every call	Yes

- **Caller-saved:** if the caller still needs r0–r3 or r12 after the call, it saves them itself;
- **Callee-saved:** a function that wants to modify r4 must restore it before returning.